

```
#step 1 : import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# step 2 : load dataset
df = pd.read_csv('/content/random_stock_market_dataset.csv')

# step 3 : data overview
print(df.info())
print(df.head())

# step 4 : summary statistics
print(df.describe())

# step 5 : check for missing values
print(df.isnull().sum())






# step 6 : outlier detection (IQR method)
numeric = df.drop('Date', axis=1)
q1 = numeric.quantile(0.25)
q3 = numeric.quantile(0.75)
iqr = q3 - q1
outliers = ((numeric < (q1 - 1.5 * iqr)) | (numeric > (q3 + 1.5 * iqr))).sum()
print('outliers per column:\n', outliers)

# step 6 : Correlation Heatmap
import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(numeric.corr(), annot=True)
plt.title('Correlation Heatmap')
plt.show()

# step 7 : Basic Visualizations
numeric.hist(figsize=(12, 8))
plt.show()

# step 8 : Save any changes (Optional)
df.to_csv('cleaned_stock_market_data.csv', index=False)
```

```
<class 'pandas.core.frame.DataFrame'>
...
RangeIndex: 60 entries, 0 to 59
Data columns (total 6 columns):
 #   Column  Non-Null Count  Dtype  
--- 
 0   Date     60 non-null    object  
 1   Open     60 non-null    float64 
 2   High     60 non-null    float64 
 3   Low      60 non-null    float64 
 4   Close    60 non-null    float64 
 5   Volume   60 non-null    int64  
dtypes: float64(4), int64(1), object(1)
memory usage: 2.9+ KB
None
Date      Open     High     Low      Close    Volume
0 2024-01-01  296.45  307.31  293.96  303.72    93133
1 2024-01-02  198.11  193.10  187.21  191.48    64993
2 2024-01-03  197.41  208.64  193.37  205.89    78326
3 2024-01-04  253.13  262.67  248.67  258.95    17358
4 2024-01-05  241.35  253.09  238.99  252.28    28847
Open      High     Low      Close    Volume
count   60.000000  60.000000  60.000000  60.000000  60.000000
mean   310.552000  322.589833  308.056833  320.412667  95485.516667
std    108.859851  109.095813  108.986426  109.157159  54895.342146
min    112.680000  123.780000  109.430000  121.270000  13193.000000
25%   222.957500  238.247500  220.042500  237.242500  48816.500000
50%   303.240000  313.800000  302.385000  309.385000  93850.500000
75%   396.150000  415.915000  392.905000  413.462500  130917.250000
max   492.790000  501.670000  492.300000  500.420000  195189.000000
Date      0
Open      0
High     0
Low      0
Close    0
Volume   0
dtype: int64
outliers per column:
Open      0
High     0
Low      0
Close    0
Volume   0
dtype: int64
```



[2]

✓ Os

```
▶ import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import Binarizer, MinMaxScaler, LabelEncoder

# Load the dataset
file_path = '/content/random_stock_market_dataset.csv'
try:
    df = pd.read_csv(file_path)
except FileNotFoundError:
    print(f"Error: File not found at {file_path}")
    # Exit or raise error if file not found
    exit()

# Create a copy to store modifications
df_processed = df.copy()

print("--- Original Data Head ---")
print(df_processed.head())
print("\n")
```

[2]
✓ 0s

```
# --- 1. Statistical Measures (Mean, Median, Mode) ---
print("--- 1. Statistical Measures ---")
# Define which columns are numerical for these stats
numerical_cols = ['Open', 'High', 'Low', 'Close', 'Volume']

# Mean
print("Mean:")
print(df_processed[numerical_cols].mean())
print("\n")

# Median
print("Median:")
print(df_processed[numerical_cols].median())
print("\n")

# Mode
print("Mode:")
# Mode can return multiple values (e.g., if two values appear with
# the same highest frequency). We print just the first row of modes.
print(df_processed[numerical_cols].mode().iloc[0])
print("\n")
```

[2]

✓ 0s

```
# --- 2. Imputation ---
print("--- 2. Imputation (Handling Missing Values) ---")
# First, check if there are any missing values in the whole DataFrame
missing_values_count = df_processed.isnull().sum().sum()
if missing_values_count == 0:
    print("No missing values found in the dataset. Imputation is not necessary.")
else:
    print(f"Found {missing_values_count} total missing values. Applying imputation...")
    # This example imputes with the mean for numerical columns.
    # You could choose 'median' or 'most_frequent' as other strategies.
    num_imputer = SimpleImputer(strategy='mean')

    for col in numerical_cols:
        if df_processed[col].isnull().any():
            print(f"Imputing column: {col}")
            # SimpleImputer expects 2D data, so we use [[col]]
            df_processed[col] = num_imputer.fit_transform(df_processed[[col]])

    print("Imputation applied (if any values were missing).")
print("\n")
```



```
# --- 3. Binarization ---
print("--- 3. Binarization ---")
# Binarize the 'Volume' column based on its mean
volume_mean = df_processed['Volume'].mean()
print(f"Binarizing 'Volume' column. Threshold (mean) = {volume_mean:.2f}")

# Initialize the Binarizer with the threshold
binarizer = Binarizer(threshold=volume_mean)

# Note: Binarizer (and most sklearn transformers) expects a 2D array.
# .values returns a NumPy array. .reshape(-1, 1) converts it to 2D.
df_processed['Volume_Binarized'] = binarizer.fit_transform(df_processed[['Volume']].values)

print("Binarization complete. New column 'Volume_Binarized' added.")
print(df_processed[['Volume', 'Volume_Binarized']].head())
print("\n")
```

```
▶ # --- 4. Normalization (Min-Max Scaling) ---
print("--- 4. Normalization (Min-Max Scaling) ---")
# Normalize the 'Close' column to the [0, 1] range
scaler = MinMaxScaler()
df_processed['Close_Normalized'] = scaler.fit_transform(df_processed[['Close']].values)

print("Normalization complete. New column 'Close_Normalized' added.")
print(df_processed[['Close', 'Close_Normalized']].head())
print("\n")

# --- 5. & 6. Label Encoding and One-Hot Encoding ---
print("--- 5. & 6. Label and One-Hot Encoding ---")
# These techniques are for categorical data.
# We will create a categorical column from 'Date' to demonstrate.

# Ensure 'Date' is in datetime format
df_processed['Date'] = pd.to_datetime(df_processed['Date'])

# Create a new categorical column: 'Day_of_Week'
df_processed['Day_of_Week'] = df_processed['Date'].dt.day_name()
print("Created new categorical column 'Day_of_Week':")
print(df_processed[['Date', 'Day_of_Week']].head())
print("\n")
```

[2]
✓ 0s

```
# 5. Label Encoding
print("Applying Label Encoding to 'Day_of_Week'...")
le = LabelEncoder()
df_processed['Day_of_Week_LabelEncoded'] = le.fit_transform(df_processed['Day_of_Week'])
print("Label Encoding complete. New column 'Day_of_Week_LabelEncoded' added.")

# Show the mapping created by the LabelEncoder
print("Label Encoder Mapping:")
for i, class_name in enumerate(le.classes_):
    print(f"{class_name} -> {i}")
print("\n")

# 6. One-Hot Encoding
print("Applying One-Hot Encoding to 'Day_of_Week'...")
# pd.get_dummies is a simple way to get one-hot encoding
one_hot_df = pd.get_dummies(df_processed['Day_of_Week'], prefix='Day')

# Join the new one-hot columns back to the main DataFrame
df_processed = pd.concat([df_processed, one_hot_df], axis=1)
print("One-Hot Encoding complete. New 'Day_*' columns added.")
print("\n")
```

[2]

✓ 0s

```
# --- Final Result ---
print("--- Final Processed Data (Head) ---")
# Print the head of the DataFrame to show all new columns
print(df_processed.head())
print("\n")

# Save the processed DataFrame to a new CSV file
output_file = 'stock_data_processed.csv'
df_processed.to_csv(output_file, index=False)
print(f"All processing complete. The new DataFrame has been saved to '{output_file}'")
```

--- Original Data Head ---

	Date	Open	High	Low	Close	Volume
0	2024-01-01	296.45	307.31	293.96	303.72	93133
1	2024-01-02	190.11	193.10	187.21	191.40	64993
2	2024-01-03	197.41	208.64	193.37	205.89	70326
3	2024-01-04	253.13	262.67	248.67	258.95	17358
4	2024-01-05	241.35	253.09	238.99	252.20	20847

--- 1. Statistical Measures ---

Mean:

Open 310.552000
High 322.589833
Low 308.056833
Close 320.412667
Volume 95405.516667
dtype: float64

Median:

Open 303.240
High 313.800
Low 302.385
Close 309.385
Volume 93850.500
dtype: float64

Mode:
Open 112.68
...
High 123.78
Low 109.43
Close 121.27
Volume 13193.00
Name: 0, dtype: float64

--- 2. Imputation (Handling Missing Values) ---
No missing values found in the dataset. Imputation is not necessary.

--- 3. Binarization ---
Binarizing 'Volume' column. Threshold (mean) = 95405.52
Binarization complete. New column 'Volume_Binarized' added.

Volume	Volume_Binarized
0	93133
1	64993
2	70326
3	17358
4	20847

--- 4. Normalization (Min-Max Scaling) ---

Normalization complete. New column 'Close_Normalized' added.

	Close	Close_Normalized
0	303.72	0.481208
1	191.40	0.184966
2	205.89	0.223183
3	258.95	0.363128
4	252.20	0.345325

--- 5. & 6. Label and One-Hot Encoding ---

Created new categorical column 'Day_of_Week':

	Date	Day_of_Week
0	2024-01-01	Monday
1	2024-01-02	Tuesday
2	2024-01-03	Wednesday
3	2024-01-04	Thursday
4	2024-01-05	Friday

Applying Label Encoding to 'Day_of_Week'...
Label Encoding complete. New column 'Day_of_Week_LabelEncoded' added.
...

Label Encoder Mapping:

- Friday -> 0
- Monday -> 1
- Saturday -> 2
- Sunday -> 3
- Thursday -> 4
- Tuesday -> 5
- Wednesday -> 6

Applying One-Hot Encoding to 'Day_of_Week'...
One-Hot Encoding complete. New 'Day_*' columns added.

--- Final Processed Data (Head) ---

	Date	Open	High	Low	Close	Volume	Volume_Binarized	\
0	2024-01-01	296.45	307.31	293.96	303.72	93133		0
1	2024-01-02	190.11	193.10	187.21	191.40	64993		0
2	2024-01-03	197.41	208.64	193.37	205.89	70326		0
3	2024-01-04	253.13	262.67	248.67	258.95	17358		0
4	2024-01-05	241.35	253.09	238.99	252.20	20847		0

```
... 0      0.481208    Monday          1      False  
1      0.184966    Tuesday         5      False  
2      0.223183    Wednesday       6      False  
3      0.363128    Thursday        4      False  
4      0.345325    Friday          0      True  
  
Day_Monday  Day_Saturday  Day_Sunday  Day_Thursday  Day_Tuesday  \\\n0      True        False        False        False        False  
1      False       False        False        False        True  
2      False       False        False        False        False  
3      False       False        False        True         False  
4      False       False        False        False        False  
  
Day_Wednesday  
0      False  
1      False  
2      True  
3      False  
4      False
```

All processing complete. The new DataFrame has been saved to 'stock_data_processed.csv'